

0Final Project

ECGR 6181/8181 Embedded Operating Systems

A priority based pre-threaded image processing server

The goal of the project is to develop an image processing server that interacts with clients through a pool of worker threads using the producer consumer model. The server consists of a main thread and a set of worker threads with the main thread running at a higher priority than the worker threads. The main thread repeatedly accepts connection requests from clients and places the resulting connected descriptors in a bounded buffer. Each worker thread repeatedly removes a descriptor from the buffer, services the client, and then waits for the next descriptor. All the threads are scheduled with the FIFO policy. The image server applies image processing operations on the image input by the client. These image processing operations are derived from the open source OpenCV library.

For this project, you can use the following C/C++ source code available on the shared Google Drive link posted on Canvas.

1. Multi-threaded producer-consumer: An array is used to implement a bounded buffer. Producer and consumer threads write and read data from this buffer synchronizing with locks and condition variables using the Pthread API
2. Real time scheduling and priorities in Linux: Demonstrates how to assign real time priority and scheduling policy to threads on Linux. Note that this code needs root permission to run (sudo in Ubuntu).
3. Delay: Spins in a loop for delay seconds. Useful to emulate the execution of long-running threads.

The project has the following milestones

1. Install OpenCV and follow the tutorial for converting a color image to greyscale.
http://docs.opencv.org/3.0-last-rst/doc/tutorials/introduction/table_of_content_introduction/table_of_content_introduction.html
http://docs.opencv.org/3.0-last-rst/doc/tutorials/introduction/load_save_image/load_save_image.html#load-save-image
2. Using the echo client & server presented in class as a template, develop the image processing server and client. The client inputs the user-input color image to the server. The server listens to connection requests at any unused port. The client connects to the server on the loopback IP address. The server should process the image converting to greyscale and send it back to the client. The client displays both the color and greyscale image. Note that the client and server are different processes running on the same machine.
3. Write a threaded version of the image processing server. On every connection request, the server spawns a new thread to process the request. Note that each connection requires a separate connection descriptor (connfd) to prevent races. Use malloc to dynamically allocate memory for the individual connection descriptors.

4. Write a pre-threaded version of the image processing server. Using a thread for each connection request is expensive due to the overheads involved in thread creation and destruction. A better approach is to have a pre-threaded server where a fixed number of worker threads (thread pool) process the connection requests. The main thread (master) that accepts connection requests puts the connection descriptors in a bounded buffer and the individual worker threads remove requests from the bounded buffer before processing.
5. Extend the pre-threaded image processing server so that the master thread has higher priority than the worker threads.

Demo: Set the number of worker threads to the one less than the number of cores on your system. Set the buffer size to twice the number of threads. Insert a delay of 30 seconds for the worker threads. Now send a series of different client requests using multiple clients. Since the master is running at a higher priority and the worker threads are slow to process the requests (due to the delay), the master will preempt the workers whenever there is a request from the client, resulting in the master finding the connection descriptor buffer full. Print appropriate messages from the master and worker threads to demonstrate that the master thread is indeed running at a higher priority.

Please organize the different parts of your code in separate C/C++ files and use the gnu make utility in compiling your code. Use of a version control system such as Git to manage your project is highly recommended (<https://git-scm.com/book/en/v2/Getting-Started-About-Version-Control>).

If your final milestone does not work, be prepared to demonstrate milestone 2, 3, and 4 for partial credit.

Implement at least 3 additional image processing functions. The client should be able to specify what kind of processing it needs as a part of the input.

See the following reference. You can use other library routines as well.

http://docs.opencv.org/2.4/doc/tutorials/imgproc/table_of_content_imgproc/table_of_content_imgproc.html